

Table of Contents

Version.....	3
Why Apache Kafka	4
Apache Kafka use cases.....	6
Kafka core API.....	7
Kafka extended API.....	8
Confluent components Schema Registry and REST Proxy	9
Administration and Monitoring tools.....	10
Course series.....	11
Kafka in Enterprise Architecture	12
Quiz.....	13
How to ask question	17
Apache Kafka core concepts	18
Topics and Partitions.....	18
Brokers.....	20
Brokers and Topics.....	21
Topic replication factor.....	22
Concept of Leader for a partition.....	24
Producers.....	25
Producers: Message keys	27
Consumers	28
Consumer groups	29
Consumer offsets	30
Zookeeper.....	31
Kafka guarantees	33
Delivery semantics for consumers	34
Quiz.....	35
Install Docker and Docker compose.....	47
Install Docker	47
Kafka Docker for development	47

Start Kafka	47
Apache Kafka: Hands-on practice	52
Topic operations: create, list, delete, describe.....	52
Kafka cli tool	52
docker run -it --rm --net=host landoop/fast-data-dev bash	52
kafka-topics --zookeeper localhost:2181 --list	53
kafka-topics --zookeeper localhost:2181 --create --topic blue_topic --partitions 3 --replication-factor 1	53
kafka-topics --zookeeper localhost:2181 --list	55
kafka-topics --zookeeper localhost:2181 --describe --topic blue_topic	56
kafka-topics --zookeeper localhost:2181 --delete --topic blue_topic	56
Another topic test	57
kaftopics --zookeeper localhost:2181 --create --topic blue_topic3 --partitions 3 --replication-factor 1	57
kafka-console-producer --broker-list localhost:9092 --topic blue_topic3	57
Comparison of the topic data	60
Consume topic	61
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning	61
.....	61
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 0	62
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 1	62
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 2	62
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --consumer-property group.id=blgrp1	63
Kafka Topics UI.....	64
Writing your own producer.....	65
Kafka API	65
Akka Streams	66

Apache Kafka

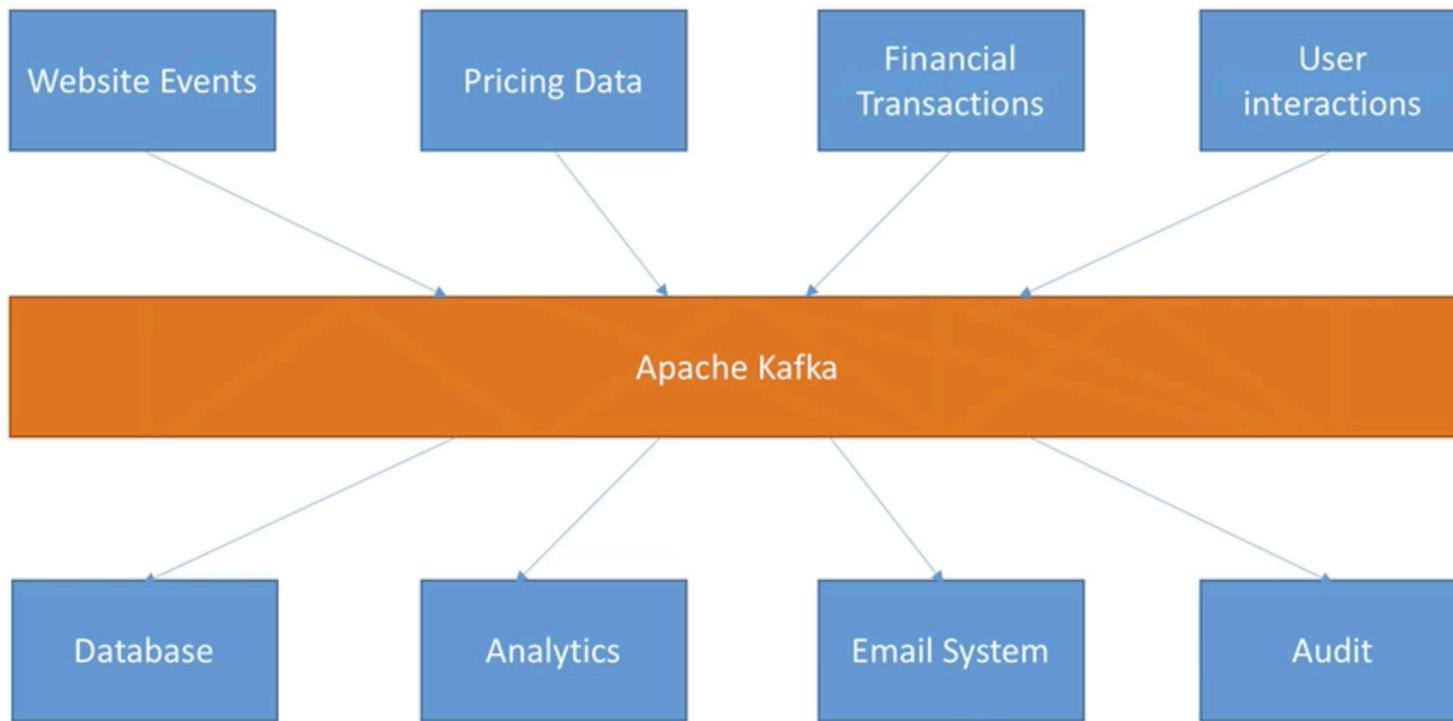
Version

Author	Version	Last update	Comment
luwenwu@cn.ibm.com	V1.0	04/11/18	Create the document

Introduction to Kafka and the ecosystem

Why Apache Kafka

Why Apache Kafka: Decoupling of data streams





Why Apache Kafka

- Distributed, resilient architecture, fault tolerant
- Horizontal scalability
- High performance (latency of less than 10ms) – real time
- Used by the 2000+ firms:
 - LinkedIn
 - Netflix
 - AirBnB
 - Yahoo
 - Walmart

Apache Kafka use cases

Apache Kafka: Use cases

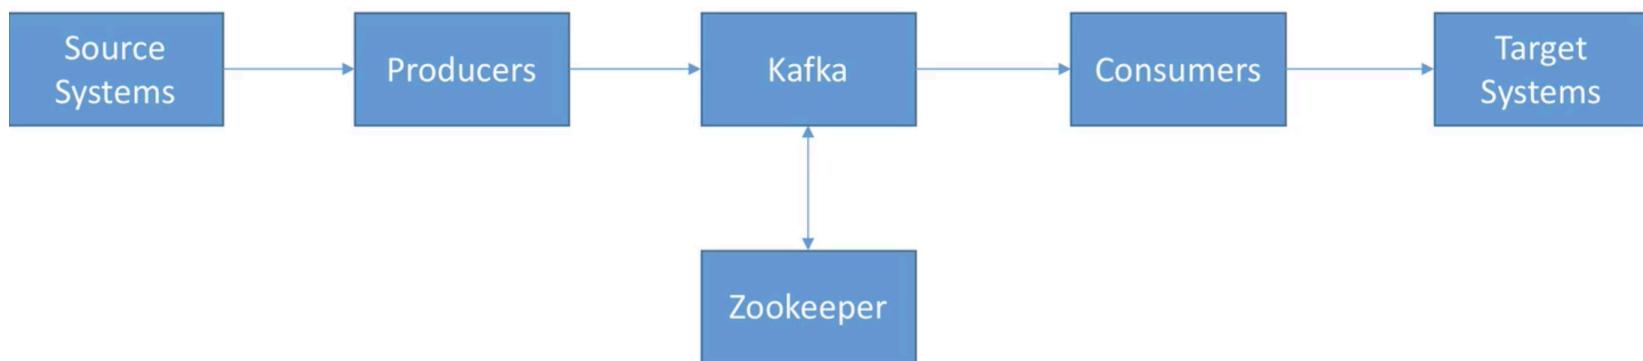


- Messaging System
- Activity Tracking
- Gather metrics from many different locations
- Application Logs gathering
- Stream processing (with the Kafka Streams API or Spark for example)
- De-coupling of system dependencies
- Integration with Spark, Flink, Storm, Hadoop, and many other Big Data technologies

Kafka core API

Kafka Ecosystem

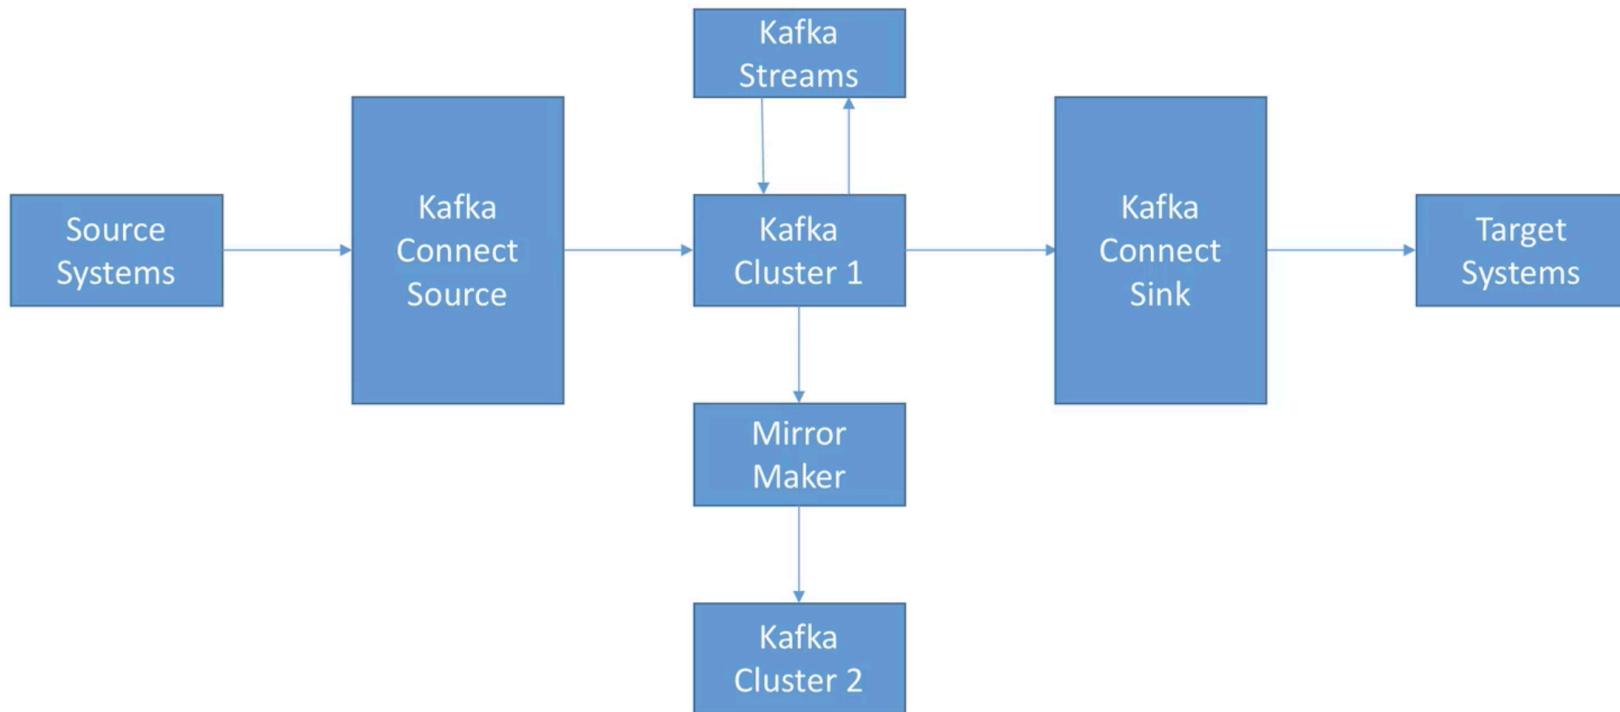
Kafka Core



Kafka extended API

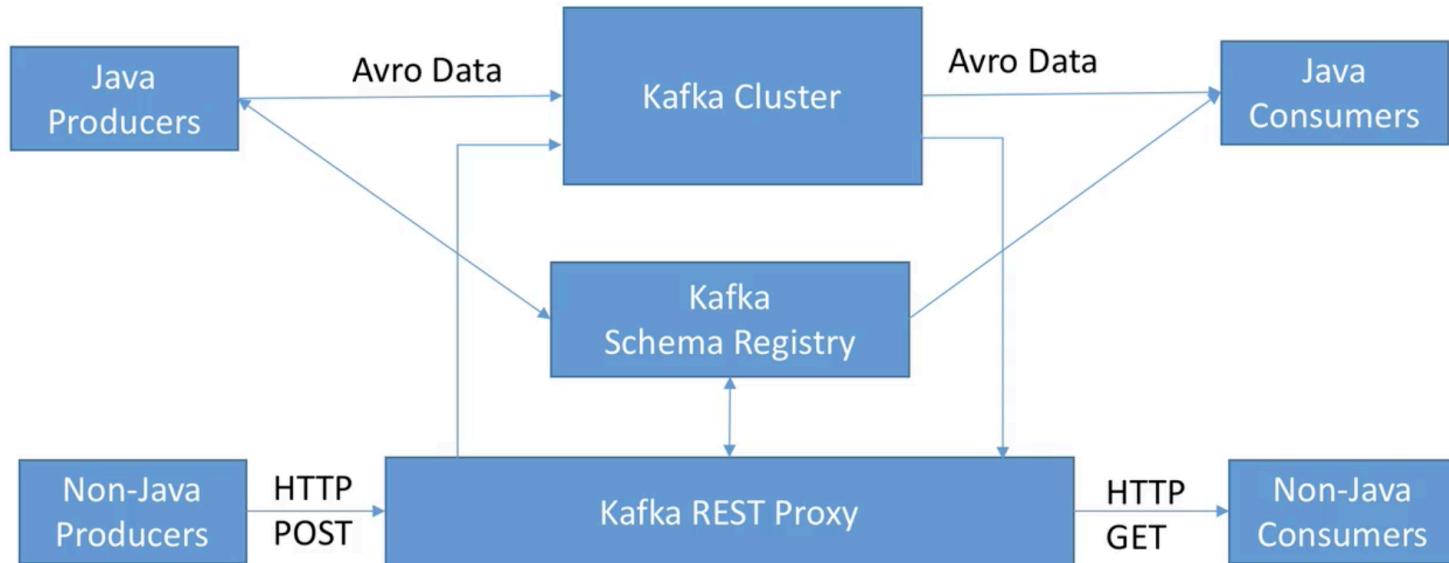
Kafka Ecosystem

Kafka Extended API



Confluent components Schema Registry and REST Proxy

Kafka Ecosystem: Confluent Components Schema Registry and REST Proxy



Administration and Monitoring tools

Kafka Ecosystem: Administration and Monitoring Tools



- Topics UI (Landoop): View the topics content
- Schema UI (Landoop): Explore the Schema registry
- Connect UI (Landoop): Create and monitor Connect tasks
- Kafka Manager (Yahoo): Overall Kafka Cluster Management
- Burrow (LinkedIn): Kafka Consumer Lag Checking
- Exhibitor (Netflix): Zookeeper Configuration, Monitoring, Backup
- Kafka Monitor (LinkedIn): Cluster health monitoring
- Kafka Tools (LinkedIn): Broker and topics administration tasks simplified
- Kafkat (Airbnb): More broker and topics administration tasks simplified
- JMX Dump: Dump JMX metrics from Brokers
- Control Centre / Auto Data Balancer / Replicator (Confluent): Paid tools

Course series



Apache Kafka Series

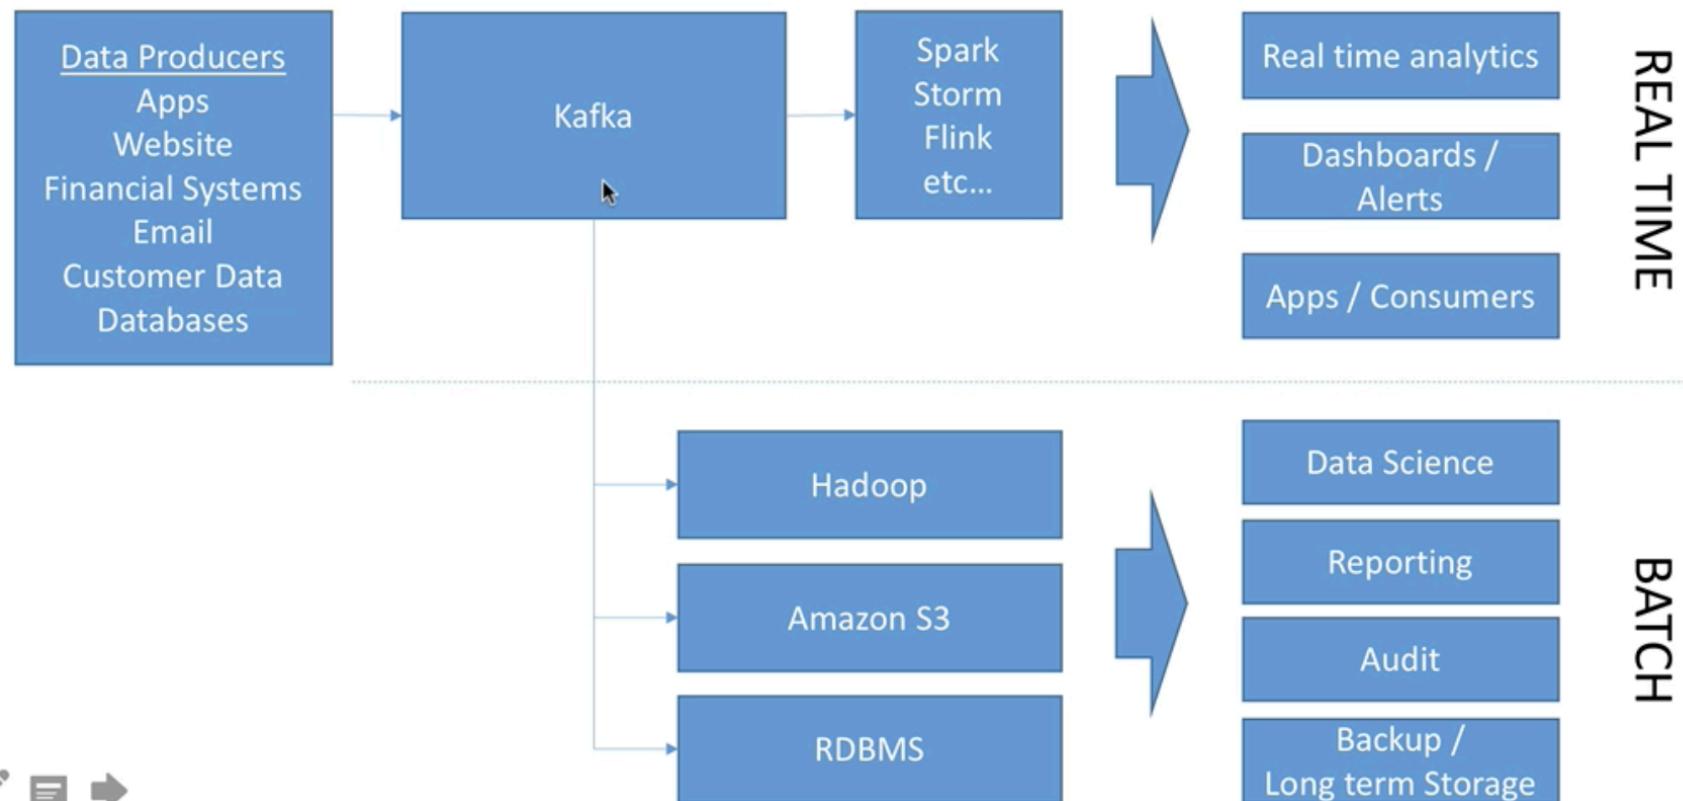
-
1. **Kafka For Beginners:** get a strong base for Kafka, basic operations, write your first producers and consumers
 2. **Kafka Connect API:** Understand how to import / export data to / from Kafka
 3. **Kafka Streams API:** Learn how to process and transform data within Kafka
 4. **Kafka Cluster Setup & Administration:** Get a deep understanding of how Kafka & Zookeeper works, how to Setup Kafka and various administration tasks in AWS.
 5. **Confluent Components:** REST Proxy and Schema Registry
 6. **Kafka Security:** Setup Kafka security in a Cluster and Integrate your applications with Kafka Security

1 for all; 2, 3 for developers; 4 for administrators

Kafka in Enterprise Architecture



Kafka in the Enterprise Architecture



Quiz

Question 1:

Kafka is...

- A new communication protocol
- A middle layer to decouple your real time data pipelines
- A database

Question 2:

Kafka is not good for

Data Streaming

Running computations such as data aggregations

De-coupling your data pipelines

Storing data in a distributed system

Question 3:

Kafka doesn't have the following capability

distributed system

horizontal scaling

fault tolerance

high performance

SQL semantics

Question 4:

Kafka is most of the time used for

real-time streaming

batch processing

How to ask question

How to Ask a Question in the Q&A

1. **Search** to see if the question has already been asked
2. **Reference** the Lecture number
3. **Take screenshots**
4. **Copy** the part of the log that has the error
5. Explain in **details** what you did, and what the error is.
6. **Be as explicit as possible**. I can't read your mind nor see your pc
7. Wait. I'm a human, and I will try to reply in time

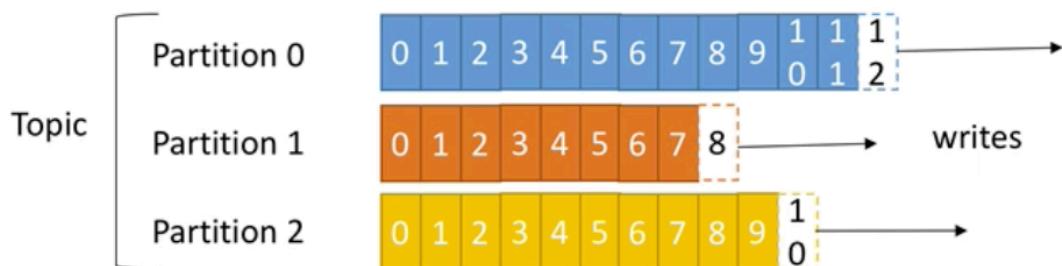
Apache Kafka core concepts

Topics and Partitions

Topics and partitions



- **Topics:** a particular stream of data
 - Similar to a table in a database (without all the constraints)
 - You can have as many topics as you want
 - A topic is identified by its name
- **Topics are split in partitions**
 - Each partition is ordered
 - Each message within a partition gets an incremental id, called offset





Topics and partitions

- Offset only have a meaning for a specific partition.
 - E.g. offset 3 in partition doesn't represent the same data as offset 3 in partition 1
- Order is guaranteed only within a partition (not across partitions)
- Data is kept only for a limited time (default is two weeks)
- Once the data is written to a partition, it can't be changed (immutability)
- Data is assigned randomly to a partition unless a key is provided (more on this later)
- You can have as many partitions per topic as you want

Brokers

Brokers



- A Kafka cluster is composed of multiple brokers (servers)
- Each broker is identified with its ID (integer)
- Each broker contains certain topic partitions
- After connecting to any broker (called a bootstrap broker), you will be connected to the entire cluster
- A good number to get started is 3 brokers, but some big clusters have over 100 brokers

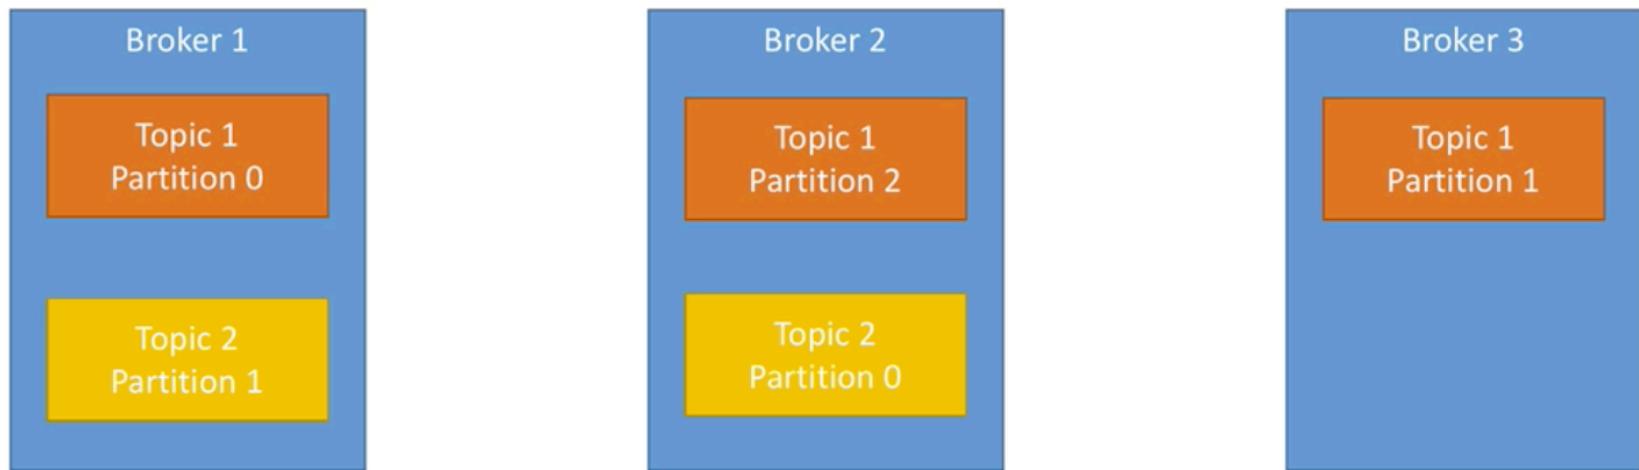


Brokers and Topics



Brokers and topics

- Example of 2 topics (3 partitions and 2 partitions)



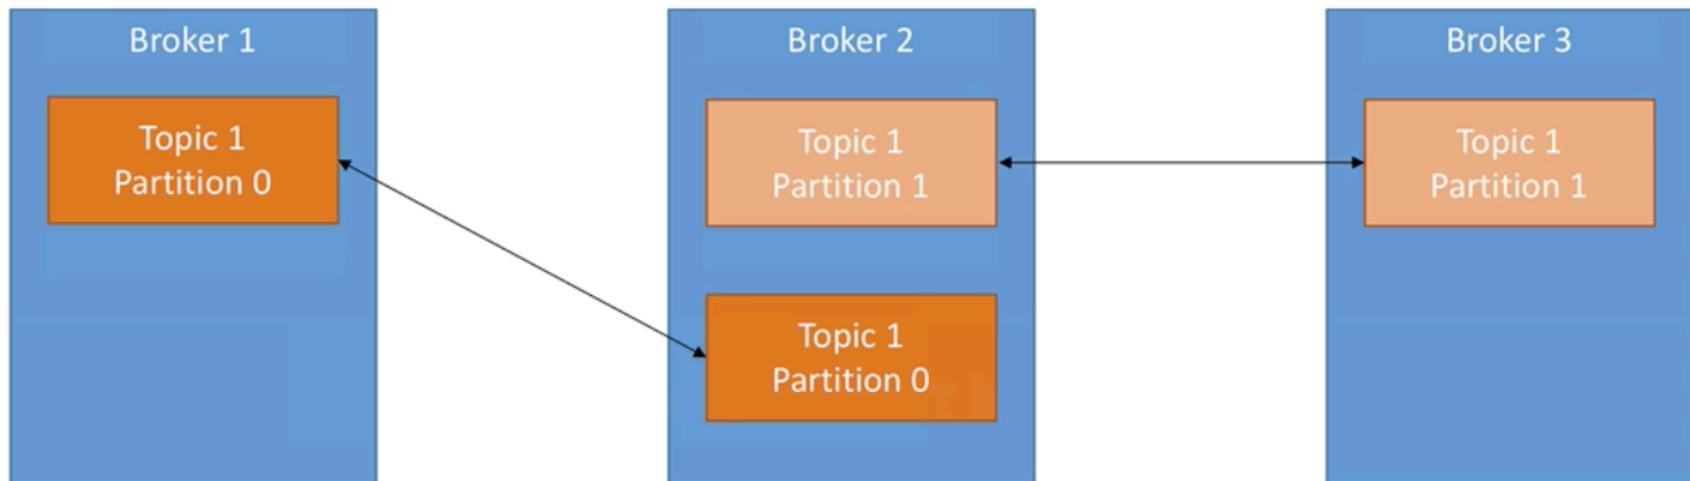
- Data is distributed and Broker 3 doesn't have any Topic 2 data

Topic replication factor

Topic replication factor



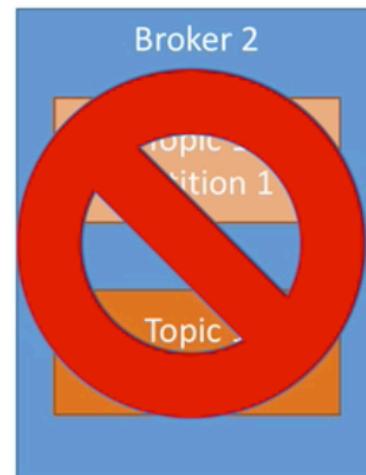
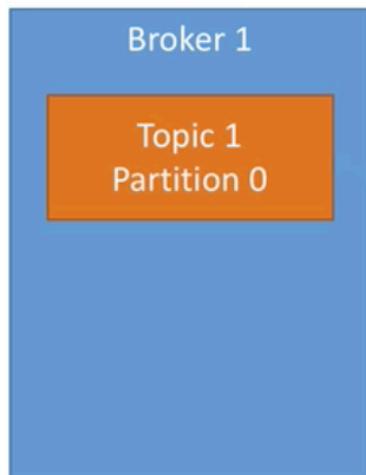
- Topics should have a replication factor > 1 (usually between 2 and 3)
- This way if a broker is down, another broker can serve the data
- Example: Topic with 2 partitions and replication factor of 2





Topic replication factor

- Example: we lost Broker 2
- Result: Broker 1 and 3 can still serve the data

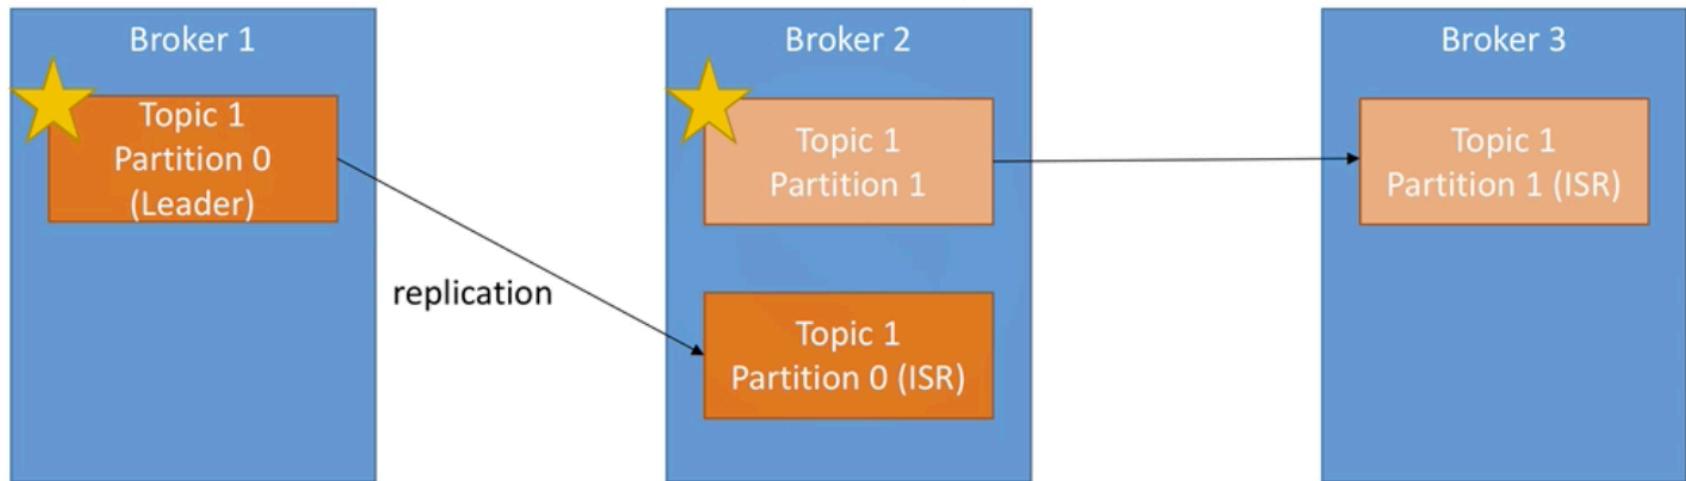


Concept of Leader for a partition



Concept of Leader for a partition

- **At any time only 1 broker can be a leader for a given partition**
- **Only that leader can receive and serve data for a partition**
- The other brokers will synchronize the data
- There each partition has: one leader, and multiple ISR (in-sync replica)

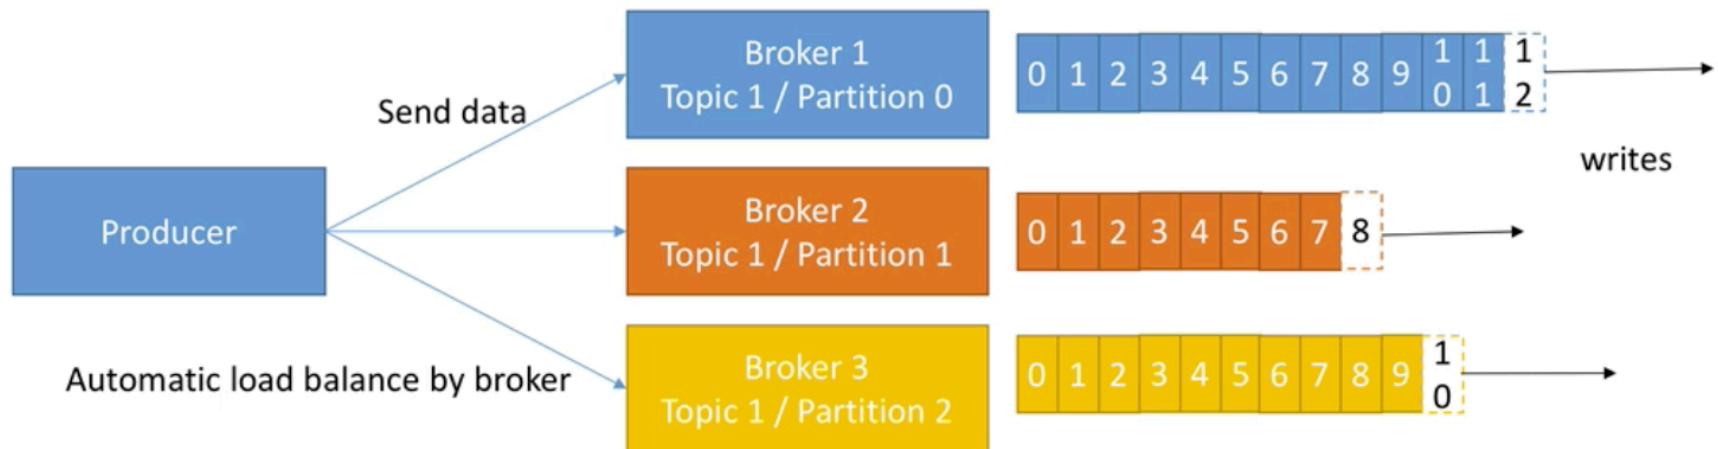


Producers

Producers



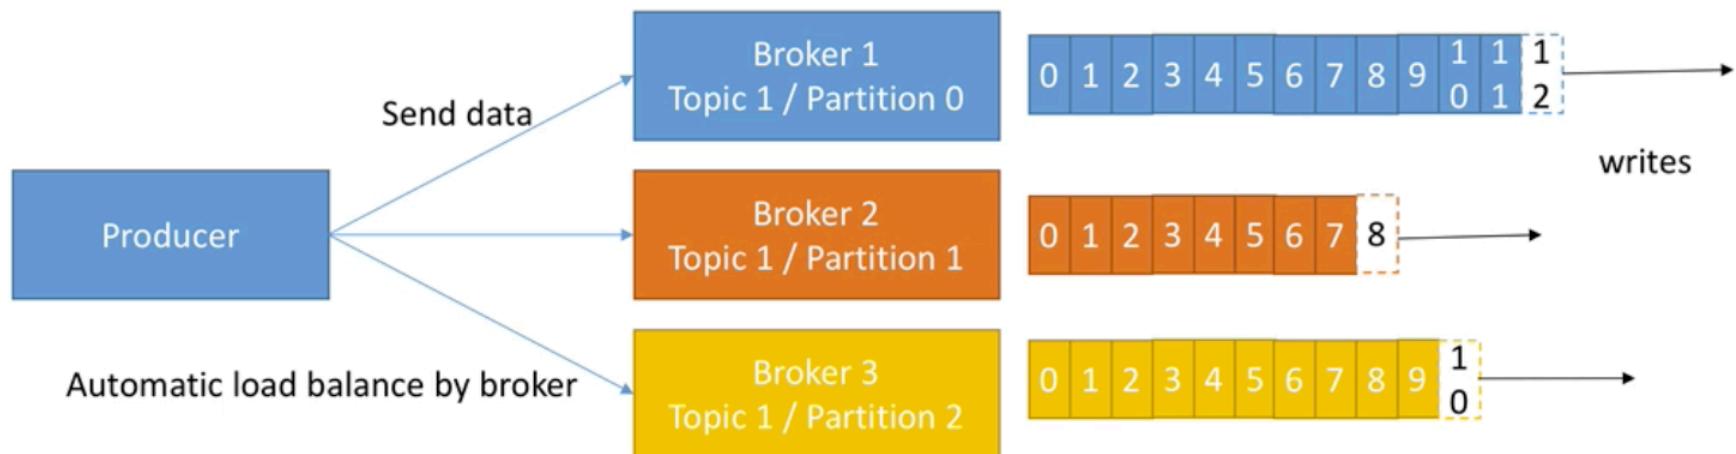
- Producers write data to topics.
- **They only have to specify the topic name and one broker to connect to, and Kafka will automatically take care of routing the data to the right brokers**





Producers

- Producers can choose to receive acknowledgment of data writes:
 - Acks=0: Producer won't wait for acknowledgment (possible data loss)
 - Acks=1: Producer will wait for leader acknowledgment (limited data loss)
 - Acks=all: Leader + replicas acknowledgment (no data loss)

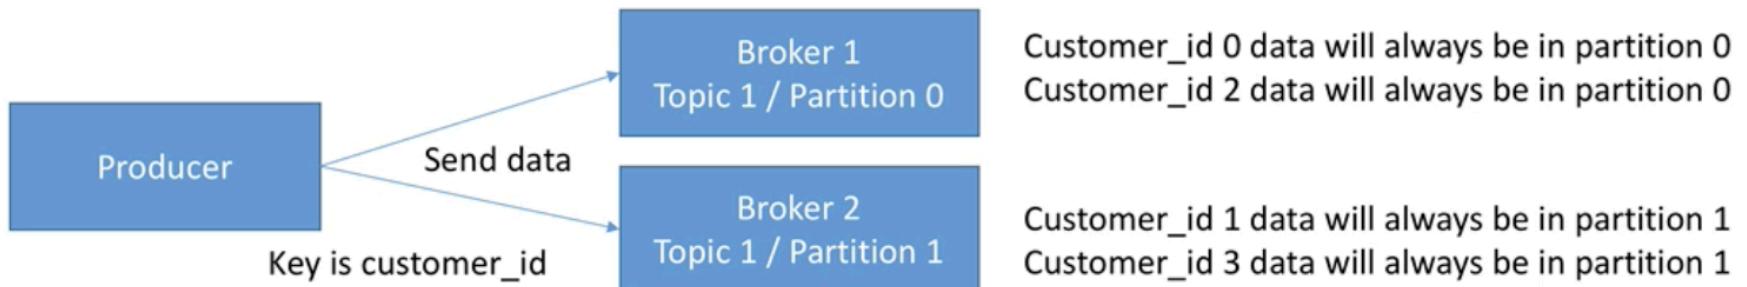


Producers: Message keys



Producers: Message keys

- Producers can choose to send a key with the message
- If a key is sent, then the producer has the guarantee that all messages for that key will always go to the same partition
- This enables to guarantee ordering for a specific key

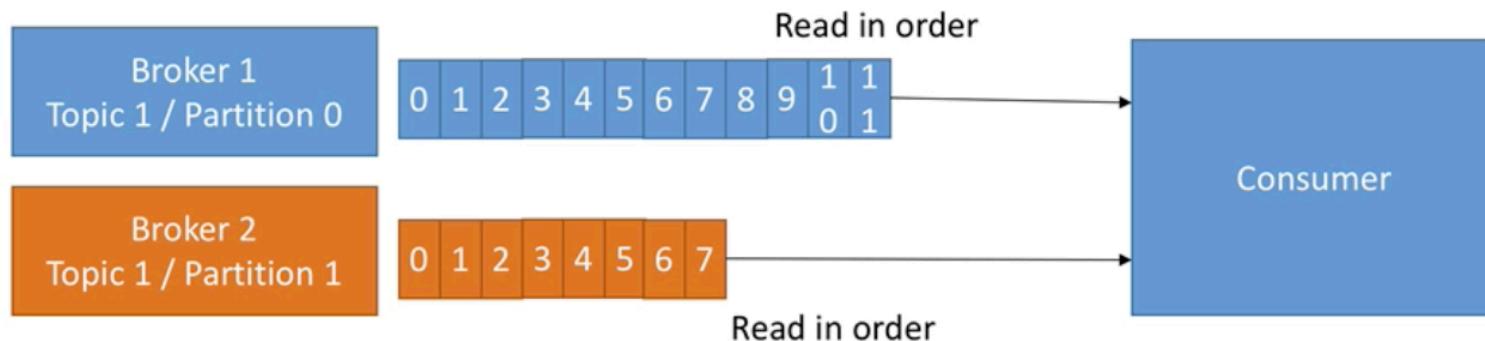


Consumers

Consumers



- Consumers read data from a topic
- **They only have to specify the topic name and one broker to connect to, and Kafka will automatically take care of pulling the data from the right brokers**
- Data is read in order **for each partitions**

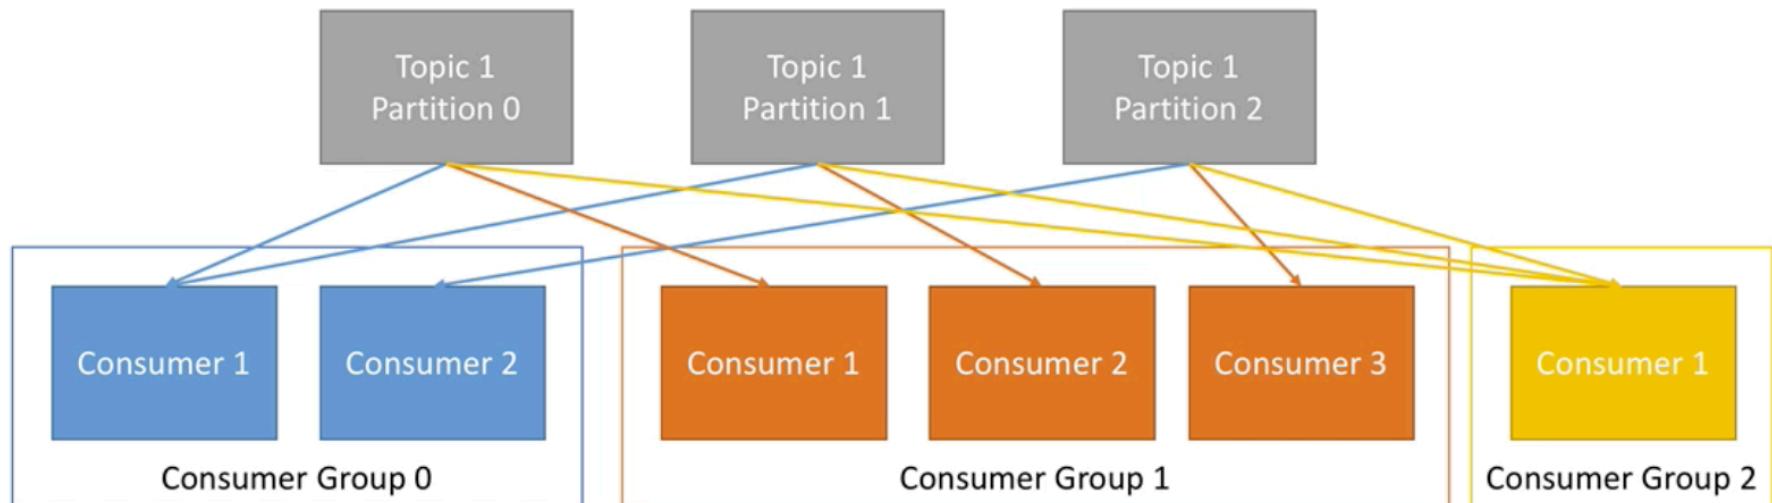


Consumer groups

Consumer Groups



- Consumers read data in consumer groups
- Each consumer within a group reads from exclusive partitions
- You cannot have more consumers than partitions (otherwise some will be inactive)

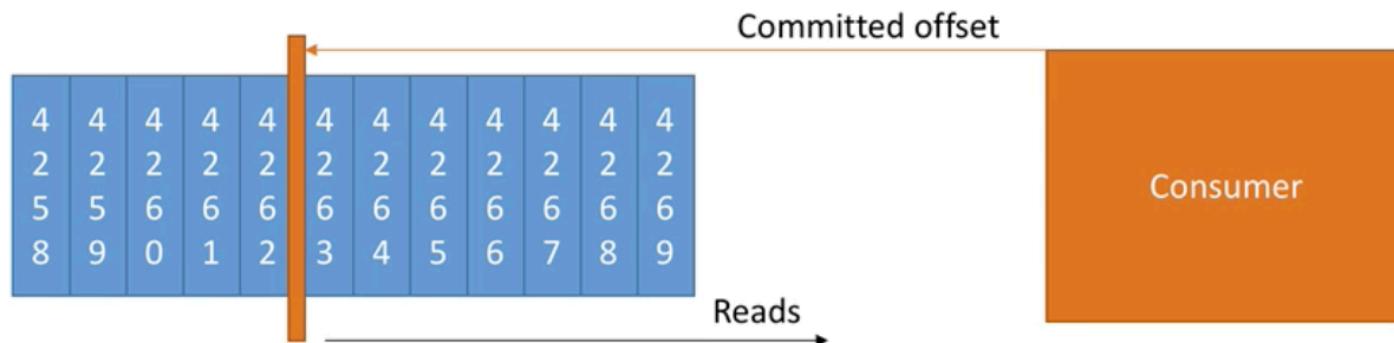


Consumer offsets



Consumer Offsets

- Kafka stores the offsets at which a consumer group has been reading
- The offsets commit live in a Kafka topic named “`__consumer_offsets`”
- When a consumer has processed data received some Kafka, it should be committing the offsets
- If a consumer process dies, it will be able to read back from where it left off thanks to consumer offsets!



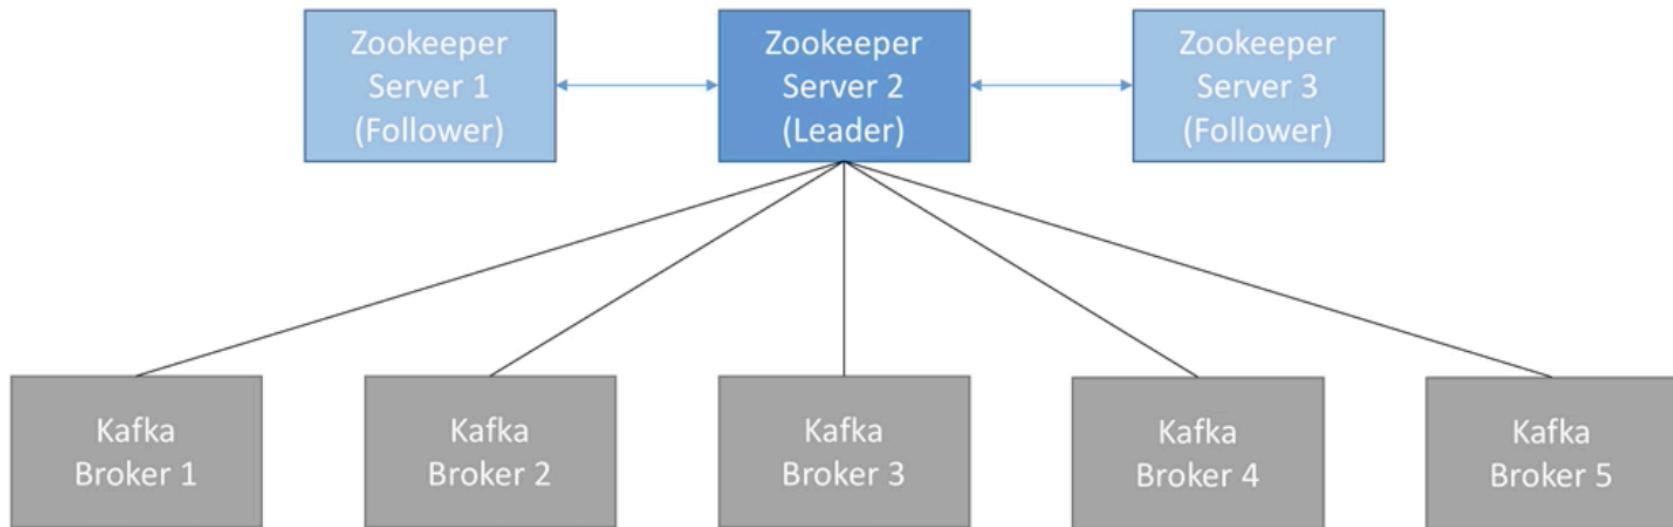
Zookeeper

Zookeeper



- Zookeeper manages brokers (keeps a list of them)
- Zookeeper helps in performing leader election for partitions
- Zookeeper sends notifications to Kafka in case of changes (e.g. new topic, broker dies, broker comes up, delete topics, etc....)
- **Kafka can't work without Zookeeper**
- Zookeeper usually operates in an odd quorum (cluster) of servers (3, 5, 7)
- Zookeeper has a leader, the rest of the servers are followers

Zookeeper



Kafka guarantees

Kafka Guarantees



- Messages are appended to a topic-partition in the order they are sent
- Consumers read messages in the order stored in a topic-partition
- With a replication factor of N, producers and consumers can tolerate up to N-1 brokers being down
- This is why a replication factor of 3 is a good idea:
 - Allows for one broker to be taken down for maintenance
 - Allows for another broker to be taken down unexpectedly
- As long as the number of partitions remains constant for a topic (no new partitions), the same key will always go to the same partition

Delivery semantics for consumers



Delivery semantics for consumers

- As learned before, consumers choose when to commit offsets.
- At most once: offsets are committed as soon as the message is received. If the processing goes wrong, the message will be lost (it won't be read again).
- At least once: offsets are committed after the message is processed. If the processing goes wrong, the message will be read again. This can result in duplicate processing of messages. Make sure your processing is **idempotent** (i.e. processing again the messages won't impact your systems)
- Exactly once: Very difficult to achieve / needs strong engineering.

Bottom line: most often you should use at least once processing and ensure your transformations / processing are idempotent

Quiz

Question 1:

Kafka topics...

always have 1 partition

can have as many partitions as desired

Question 2:

Offsets only at the level of

the topic

the topic-partition

Question 3:

Once sent to a topic, a message can be modified

true

false

Question 4:

Brokers are identified by

A name (string)

An ID (number)

Question 5:

Every brokers

contain all the topics and all the partitions

contain only a subset of the topics and the partitions

Question 6:

If a topic has a replication factor of 3

Each partition will live on 3 different brokers

Each partition will live on 2 different brokers

Each partition will live on 4 different brokers

Question 7:

If a topic has a replication factor of 3, what maximum number of brokers can be stopped without impacting the cluster?

1

2

3

Question 8:

Each partition can only have 1 leader, and multiple replicas

true

false

✓ Good job!

Very important: you only need to connect to one broker (any broker) and just provide the topic name you want to write to. Kafka will route your data to the appropriate brokers and partitions for you!

Question 9:

To produce data to a topic, a producer must provide the Kafka client with...

any broker from the cluster and the topic name

any broker from the cluster and the topic name and the partitions list

all the brokers from the cluster and the topic name

the list of brokers that have the data, the topic name, and the partitions list

Question 10:

To get acknowledgement of writes to only the leader, we need to use the config...

acks=1

acks=0

acks=all

Question 11:

To read data from a topic, the following configuration is needed for the consumers

- any brokers to connect to, and the topic name

- all brokers of the cluster, and the topic name

- any brokers, and the list of topic partitions

Question 12:

Two consumers that have the same group.id (consumer group id) will read from mutually exclusive partitions

true

false

Question 13:

As of Kafka Version >= 0.9.0

Zookeeper is not used

to perform a leader election

to maintain the list of all the brokers in the cluster

to send notifications to Kafka in case of any changes

to maintain consumer offsets

Install Docker and Docker compose

Install Docker

I have installed Docker and Docker compose on my Linux server. So I'd skip this step. But I'd show the repo here for reference purpose:

```
[root@sbybz3137 ~]# cat /etc/yum.repos.d/docker-ce.repo
[docker-ce-stable]
name=Docker CE Stable - $basearch
baseurl=https://download.docker.com/linux/centos/7/$basearch/stable
enabled=1
gpgcheck=1
gpgkey=https://download.docker.com/linux/centos/gpg
```

Kafka Docker for development

Kafka, Zookeeper, Schema Registry, Kafka-Connect, Landoop Tools, 20+ connectors
<https://github.com/Landoop/fast-data-dev>

Start Kafka

```
docker run --rm --net=host -e ADV_HOST=sbybz3137.sby.ibm.com landoop/fast-data-dev
```

```
[root@sbybz3137 ~]# docker run --rm --net=host landoop/fast-data-dev
Unable to find image 'landoop/fast-data-dev:latest' locally
latest: Pulling from landoop/fast-data-dev
ff3a5c916c92: Already exists
0a6f3946d366: Pull complete
38cbac368ac8: Pull complete
940e9a6559b3: Pull complete
2533ff68afb1: Pull complete
```

5310643b7010: Pull complete
4a0bf0f7b38c: Pull complete
256a5053f0c5: Pull complete
df59bd94bdd4: Pull complete
6aac547c4d90: Pull complete
1818c7d8b194: Pull complete
8297d3821b7b: Pull complete
27f5c49688ee: Pull complete
3a13574f4ddd: Pull complete
13d13524515b: Pull complete
dd4dddc9dc97: Pull complete
633bee4809e9: Pull complete
a6842c7eb2cf: Pull complete
22ae19fc8a44: Pull complete
c58290624b9c: Pull complete
09afeafbccaa3: Pull complete
30033ace6359: Pull complete
14951d7d8ded: Pull complete
6797a96a7092: Pull complete
5e958485d959: Pull complete
6e77d9fc1e81: Pull complete
99046468067f: Pull complete
def900f2c3f5: Pull complete
e3298ece27a7: Pull complete
46b35ea8b707: Pull complete
498af3d85513: Pull complete
c29a64ba4098: Pull complete
d5267660efab: Pull complete
068d760d616b: Pull complete
12c2d93a4a0c: Pull complete

Digest: sha256:041358122a185144261c95cb8abfb38d47461d585e1672ed58d3ce9800634629

Status: Downloaded newer image for landoop/fast-data-dev:latest

Starting services.

This is landoop's fast-data-dev. Kafka 0.11.0.1, Confluent OSS 3.3.1.

You may visit <http://localhost:3030> in about a minute.

2018-04-12 20:57:06,795 CRIT Supervisor running as root (no user in config file)

2018-04-12 20:57:06,795 INFO Included extra file "/etc/supervisord.d/01-fast-data.conf" during parsing

2018-04-12 20:57:06,795 INFO Included extra file "/etc/supervisord.d/99-supervisord-sample-data.conf" during parsing

2018-04-12 20:57:06,797 INFO supervisord started with pid 7

2018-04-12 20:57:07,800 INFO spawned: 'sample-data' with pid 99

2018-04-12 20:57:07,802 INFO spawned: 'zookeeper' with pid 100

2018-04-12 20:57:07,805 INFO spawned: 'caddy' with pid 101

2018-04-12 20:57:07,808 INFO spawned: 'broker' with pid 102

2018-04-12 20:57:07,811 INFO spawned: 'smoke-tests' with pid 103

2018-04-12 20:57:07,817 INFO spawned: 'connect-distributed' with pid 106

2018-04-12 20:57:07,819 INFO spawned: 'logs-to-kafka' with pid 110

2018-04-12 20:57:07,821 INFO spawned: 'schema-registry' with pid 111

2018-04-12 20:57:07,823 INFO spawned: 'rest-proxy' with pid 116

2018-04-12 20:57:09,464 INFO success: sample-data entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: zookeeper entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: caddy entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: broker entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: smoke-tests entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: connect-distributed entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: logs-to-kafka entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: schema-registry entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:57:09,464 INFO success: rest-proxy entered RUNNING state, process has stayed up for > than 1 seconds (startsecs)

2018-04-12 20:59:32,743 INFO exited: sample-data (exit status 0; expected)

2018-04-12 21:00:08,651 INFO exited: logs-to-kafka (exit status 0; expected)

2018-04-12 21:02:12,065 INFO exited: smoke-tests (exit status 0; expected)

Or other options:

If you want to have the services remotely accessible, then you need to pass in your machine's IP address or hostname that other machines can use to access it:

```
docker run --rm --net=host landoop/fast-data-dev
```

```
docker run --rm --net=host -e ADV_HOST=<IP> landoop/fast-data-dev
```

```
docker run --rm --net=host -e ADV_HOST=sbybz3137.sby.ibm.com landoop/fast-data-dev
```

<http://sbybz3137.sby.ibm.com:3030/>

 Kafka Development Environment
docker container powered by Landoop

SCHEMAS	TOPICS	CONNECTORS	BROKERS
15	12	1	1
SCHEMA REGISTRY UI manage avro schemas ENTER	KAFKA TOPICS UI browse topics and data ENTER	KAFKA CONNECT UI setup & manage connectors ENTER	LENSES management and monitoring ENTER

 RUNNING SERVICES
fast-data-dev

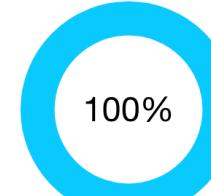
- ✓ Kafka 0.11.0.1 — Confluent OSS 3.3.1
1x Broker, 1x Schema Registry, 1x Connect Distributed Worker, 1x REST Proxy, 1x Zookeeper
- ✓ Landoop Stream Reactor 0.4.0
Source & Sink connectors collection (25+) supporting KCQL
- ✓ Landoop Schema Registry UI v0.9.4
Create, view, search, edit, validate, evolve & configure Avro schemas
- ✓ Landoop Kafka Topics UI v0.9.3
Browse and search topics, inspect data, metadata and configuration
- ✓ Landoop Kafka Connect UI v0.9.4
View, create, update and manage connectors

SERVICES PORTS

9092 : Kafka Broker	9581 : JMX
8081 : Schema Registry	9582 : JMX
8082 : Kafka REST Proxy	9583 : JMX
8083 : Kafka Connect Distributed	9584 : JMX
2181 : ZooKeeper	9585 : JMX
3030 : Web Server	

BROWSE
[- running services log files](#)

COYOTE HEALTH CHECKS
Your set up is being verified using the Coyote integration testing tool and generates working examples.



100%

Passed: 42 | Failed: 0

[VIEW RESULTS](#)

Find out more about Coyote awesome tool [here](#).

FAST DATA NEWS

- **Stream Reactor 0.4.0 and 1.0.0**
Landoop Stream Reactor hit versions 0.4.0 (targeting Kafka 0.11.0) and 1.0.0 (targeting Kafka 1.0). This is our biggest release ever with 25 connectors in total. [Find here our release notes and builds](#).
- **Schema Registry UI and Kafka Connect UI reach 0.9.4**
A twin release by our frontend team! Schema Registry UI now supports schema deletion. Kafka Connect UI adds new connectors and even better support for Kafka 0.11.0 and 1.0. Find out more in the [Schema Registry UI 0.9.4 release notes](#) and the [Kafka Connect UI 0.9.4 release notes](#).
- **Datamountaineer and Stream Reactor Join Landoop!**

Apache Kafka: Hands-on practice

Topic operations: create, list, delete, describe

Kafka cli tool

docker run -it --rm --net=host landoop/fast-data-dev bash

```
[root@sbybz3137 ~]# docker run -it --rm --net=host landoop/fast-data-dev bash
root@fast-data-dev / $ kafka-
kafka-acls                                     kafka-console-producer          kafka-preferred-replica-election  kafka-rest-start                kafka-simple-consumer-shell
kafka-avro-console-consumer                   kafka-consumer-groups           kafka-producer-perf-test         kafka-rest-stop                 kafka-streams-application-reset
kafka-avro-console-producer                   kafka-consumer-offset-checker   kafka-reassign-partitions       kafka-rest-stop-service        kafka-topics
kafka-broker-api-versions                    kafka-consumer-perf-test       kafka-replay-log-producer      kafka-run-class                kafka-verifiable-consumer
kafka-configs                                    kafka-delete-records          kafka-replica-verification    kafka-server-start             kafka-verifiable-producer
kafka-console-consumer                      kafka-mirror-maker            kafka-rest-run-class          kafka-server-stop             
```

```
kafka-topics --zookeeper localhost:2181 --list
```

```
root@fast-data-dev ~ $ kafka-topics --zookeeper localhost:2181 --list
__consumer_offsets
_schemas
backblaze_smart
connect-configs
connect-offsets
connect-statuses
coyote_test_avro
coyote_test_binary
coyote_test_json
nyc_yellow_taxi_trip_data
reddit_posts
sea_vessel_position_reports
telecom_italia_data
telecom_italia_grid
```

```
kafka-topics --zookeeper localhost:2181 --create --topic blue_topic --partitions 3 --replication-factor 1
```

WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.

Created topic "blue_topic".

← → ⌂ ⓘ sbybz3137.sby.ibm.com:3030/kafka-topics-ui/#/

KAFKA TOPICS

10 TOPICS System Topics

Search topics:

Topic	Partitions × Replication	Configurations	Format	
backblaze_smart	2 Partitions	× 1 Replication	3configs	...
blue_topic	3 Partitions	× 1 Replication		...
coyote_test_avro	1 Partitions	× 1 Replication		...
coyote_test_binary	1 Partitions	× 1 Replication		...
coyote_test_json	1 Partitions	× 1 Replication		...
nyc_yellow_taxi_trip_data	1 Partitions	× 1 Replication	3configs	...
reddit_posts	5 Partitions	× 1 Replication	3configs	avro
sea_vessel_position_reports	3 Partitions	× 1 Replication	3configs	avro
telecom_italia_data	4 Partitions	× 1 Replication	3configs	avro
telecom_italia_grid	1 Partitions	× 1 Replication	3configs	avro

TOTAL TOPICS 15

KAFKA REST
/api/kafka-rest-proxy

KAFKA-TOPICS-UI v0.9.3

BROKERS 1

Powered by [Landoop](#)

```
kafka-topics --zookeeper localhost:2181 --list
```

```
root@fast-data-dev / $ kafka-topics --zookeeper localhost:2181 --list
__consumer_offsets
_schemas
backblaze_smart
blue_topic
connect-configs
connect-offsets
connect-statuses
coyote_test_avro
coyote_test_binary
coyote_test_json
nyc_yellow_taxi_trip_data
reddit_posts
sea_vessel_position_reports
telecom_italia_data
telecom_italia_grid
```

```
kafka-topics --zookeeper localhost:2181 --describe --topic blue_topic
```

```
root@fast-data-dev / $ kafka-topics --zookeeper localhost:2181 --describe --topic blue_topic
Topic:blue_topic      PartitionCount:3      ReplicationFactor:1      Configs:
          Topic: blue_topic      Partition: 0      Leader: 0      Replicas: 0      Isr: 0
          Topic: blue_topic      Partition: 1      Leader: 0      Replicas: 0      Isr: 0
          Topic: blue_topic      Partition: 2      Leader: 0      Replicas: 0      Isr: 0
```

```
kafka-topics --zookeeper localhost:2181 --delete --topic blue_topic
```

```
root@fast-data-dev / $ kafka-topics --zookeeper localhost:2181 --delete --topic blue_topic
Topic blue_topic is marked for deletion.
Note: This will have no impact if delete.topic.enable is not set to true.
```

Another topic test

```
kaftopics --zookeeper localhost:2181 --create --topic blue_topic3 --partitions 3 --replication-factor 1  
kafka-console-producer --broker-list localhost:9092 --topic blue_topic3
```

```
>^Croot@fast-data-dev / $ kaftopics --zookeeper localhost:2181 --create --topic blue_topic3 --partitions 3 --replication-factor 1  
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.  
Created topic "blue_topic3".  
root@fast-data-dev / $ kafka-console-producer --broker-list localhost:9092 --topic blue_topic3  
>1.one  
>2.two  
>3.three  
>4.four  
>5.five  
>6.six  
>7.seven  
>8.eight  
>9.nine  
>10.ten
```

← → ⌂ ⓘ sbybz3137.sby.ibm.com:3030/kafka-topics-ui/#/cluster/fast-data-dev/topic/n/blue_topic3/data

KAFKA TOPICS

12 TOPICS System Topics

Search topics:

backblaze_smart 2 Partitions × 1 Replication 3configs	...
blue_topic 1 Partitions × 1 Replication	...
blue_topic2 3 Partitions × 1 Replication	...
blue_topic3 3 Partitions × 1 Replication	...
coyote_test_avro 1 Partitions × 1 Replication	...
coyote_test_binary 1 Partitions × 1 Replication	...
coyote_test_json 1 Partitions × 1 Replication	...
nyc_yellow_taxi_trip_data	...

blue_topic3

DATA Total Messages Fetched: 10. Data type: binary

filter

TOPIC	TABLE	RAW DATA	
Offset	Partition	Key	Value
0	2	ée	3.three
1	2	ée	6.six
2	2	ée	9.nine
0	1	ée	1.one
1	1	ée	4.four
2	1	ée	7.seven
3	1	ée	10.ten
0	0	ée	2.two
1	0	ée	5.five
2	0	ée	8.eight

↗ blue_topic3

DATA PARTITIONS 3 CONFIGURATION

Partitions

Number of Partitions: 3
Replication Factor: 1

Partition	Replica	Replica Broker	Is Leader	Is in-sync
0	0	0	true	true
1	0	0	true	true
2	0	0	true	true

Comparison of the topic data

Two screenshots of a Kafka topic viewer interface for "blue_topic3".

Left Screenshot (Table View):

- Header: "blue_topic3" (with a back arrow), "DATA", "Total Messages Fetched: 10. Data type: binary".
- Filter: "filter".
- Table Headers: "TOPIC", "TABLE", "RAW DATA".
- Table Rows (Offset, Partition, Key, Value):

Offset	Partition	Key	Value
0	2	ée	3.three
1	2	ée	6.six
2	2	ée	9.nine
0	1	ée	1.one
1	1	ée	4.four
2	1	ée	7.seven
3	1	ée	10.ten
0	0	ée	2.two
1	0	ée	5.five
2	0	ée	8.eight

Right Screenshot (Raw Data View):

- Header: "blue_topic3" (with a back arrow), "DATA", "Total Messages Fetched: 10. Data type: binary".
- Filter: "filter".
- Table Headers: "TOPIC", "TABLE", "RAW DATA".
- Raw Data (Message Log):


```

1 - [
2 -   {
3 -     "topic": "blue_topic3",
4 -     "key": "ée",
5 -     "value": "3.three",
6 -     "partition": 2,
7 -     "offset": 0
8 -   },
9 -   {
10 -    "topic": "blue_topic3",
11 -    "key": "ée",
12 -    "value": "6.six",
13 -    "partition": 2,
14 -    "offset": 1
15 -   },
16 -   {
17 -     "topic": "blue_topic3",
18 -     "key": "ée",
19 -     "value": "9.nine",
20 -     "partition": 2,
21 -     "offset": 2
22 -   },
23 -   {
24 -     "topic": "blue_topic3",
25 -     "key": "ée",
26 -     "value": "1.one",
27 -     "partition": 1,
28 -     "offset": 0
29 -   },
30 -   {
31 -     "topic": "blue_topic3",
32 -     "key": "ée",
33 -     "value": "4.four",
34 -     "partition": 1,
35 -     "offset": 1
36 -   },
37 -   {
38 -     "topic": "blue_topic3",
39 -     "key": "ée",
40 -     "value": "7.seven",
41 -     "partition": 1,
42 -     "offset": 2
43 -   },
44 -   {
45 -     "topic": "blue_topic3",
46 -     "key": "ée",
47 -     "value": "10.ten",
48 -     "partition": 1,
49 -     "offset": 3
50 -   },
51 -   {
52 -     "topic": "blue_topic3",
53 -     "key": "ée",
54 -     "value": "2.two",
55 -     "partition": 0,
56 -     "offset": 0
57 -   },
58 -   {
59 -     "topic": "blue_topic3",
60 -     "key": "ée",
61 -     "value": "5.five",
62 -     "partition": 0,
63 -     "offset": 1
64 -   },
65 -   {
66 -     "topic": "blue_topic3",
67 -     "key": "ée",
68 -     "value": "8.eight",
69 -     "partition": 0,
70 -     "offset": 2
71 -   }
72 - ]
      
```

Consume topic

```
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning
```

```
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning
3.three
6.six
9.nine
1.one
4.four
7.seven
10.ten
2.two
5.five
8.eight
```

```
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 0
```

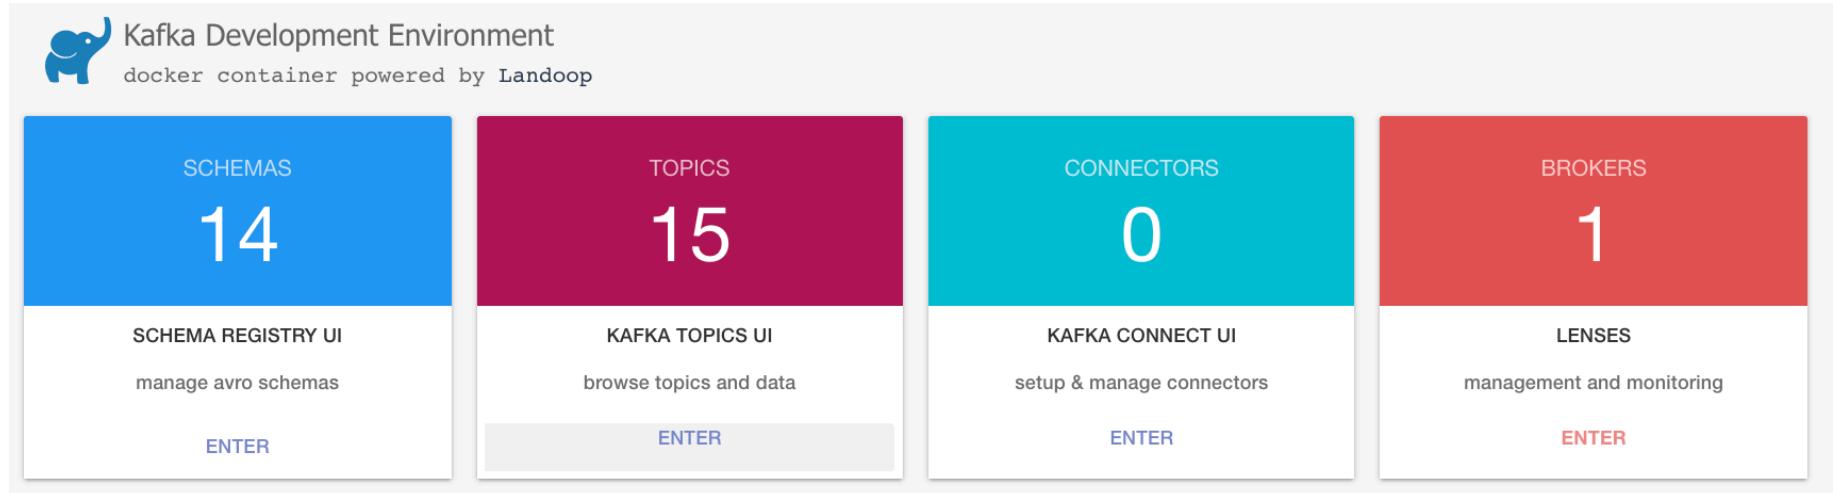
```
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 1
```

```
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 2
```

```
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 0
2.two
5.five
8.eight
^CProcessed a total of 3 messages
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 1
1.one
4.four
7.seven
10.ten
^CProcessed a total of 4 messages
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 2
3.three
6.six
9.nine
4.ten
7.seven
10.ten
2.two
5.five
8.eight
kafka-console-consumer --bootstrap-server localhost:9092 --topic blu
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --partition 2
2.two
5.five
8.eight
```

```
kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --consumer-property group.id=blgrp1
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning --consumer-property group.id=blgrp1
3.three
6.six
9.nine
1.one
4.four
7.seven
10.ten
2.two
5.five
8.eight
^CProcessed a total of 10 messages
root@fast-data-dev / $ kafka-console-consumer --bootstrap-server localhost:9092 --topic blue_topic3 --from-beginning
3.three
6.six
9.nine
1.one
4.four
7.seven
10.ten
2.two
5.five
8.eight
```

Kafka Topics UI



← → ⌂ ⓘ sbybz3137.sby.ibm.com:3030/kafka-topics-ui/#/

KAFKA TOPICS

10 TOPICS

Search topics:

- backblaze_smart
2 Partitions x 1 Replication | 3configs
- blue_topic3
3 Partitions x 1 Replication
- coyote_test_avro
1 Partitions x 1 Replication

TOTAL TOPICS

15

KAFKA REST
</api/kafka-rest-proxy>

KAFKA-TOPICS-UI v0.9.3

BROKERS

1

Writing your own producer

Kafka API

<http://kafka.apache.org/documentation/#producerapi>

To use the producer, you can use the following maven dependency:

```
1 <dependency>
2   <groupId>org.apache.kafka</groupId>
3   <artifactId>kafka-clients</artifactId>
4   <version>1.1.0</version>
5 </dependency>
6
```

To use the consumer, you can use the following maven dependency:

```
1 <dependency>
2   <groupId>org.apache.kafka</groupId>
3   <artifactId>kafka-clients</artifactId>
4   <version>1.1.0</version>
5 </dependency>
6
```

To use Kafka Streams you can use the following maven dependency:

```
1 <dependency>
2   <groupId>org.apache.kafka</groupId>
3   <artifactId>kafka-streams</artifactId>
4   <version>1.1.0</version>
5 </dependency>
6
```

Akka Streams

TODO: